

Measuring error

Ciaran Evans

Warmup

Your friend collects $n = 100$ observations with a binary response, and fits a logistic regression model:

$$Y_i \sim \text{Bernoulli}(p_i)$$
$$\log\left(\frac{p_i}{1-p_i}\right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_{99} X_{i,99}$$

Using this data, they then create a confusion matrix:

Predicted Status	Actual Status	
	negative	positive
negative	42	0
positive	0	58

What has gone wrong?

Take-aways from last class activity

- ▶ Even if the model is correctly specified, prediction performance on the **training** data is a biased estimate of prediction performance on **new** data
- ▶ The bias gets worse as the number of explanatory variables in the model increases!

Terminology and notation

Training data:

Fitted model:

New data:

Loss functions

Loss function:

Risk (error)

Training error:

Test error:

Expected test error:

What do we do??

If training error is a **biased** estimate for error on new data, how do we proceed?

Class activity

Work on the class activity (course website)