

Poisson regression

Ciaran Evans

Warmup

Work on the warmup activity (handout)

Warmup

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	0.276493	0.087294	3.167	0.00154	**
phd	0.008071	0.026358	0.306	0.75944	
ment	0.026537	0.002001	13.264	< 2e-16	***
kid5	-0.099696	0.035248	-2.828	0.00468	**

How would we interpret the coefficient on the number of articles published by the student's mentor?

Warmup

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	0.276493	0.087294	3.167	0.00154	**
phd	0.008071	0.026358	0.306	0.75944	
ment	0.026537	0.002001	13.264	< 2e-16	***
kid5	-0.099696	0.035248	-2.828	0.00468	**

What is the predicted number of articles for a student at a PhD program with a prestige score of 4, whose mentor published 30 articles in the last 3 years, and who has no children?

EDMs and GLMs

We say that a distribution belongs to the **exponential dispersion model** family if the distribution's probability function can be written in the form

$$f(y; \theta, \phi) = a(y, \phi) \exp \left\{ \frac{y\theta - \kappa(\theta)}{\phi} \right\}$$

$$f(y; \mu, \phi) = b(y, \phi) \exp \left\{ -\frac{d(y, \mu)}{2\phi} \right\}$$

For iid observations $(\mathbf{x}_1, Y_1), \dots, (\mathbf{x}_n, Y_n)$, a **generalized linear model** (GLM) consists of two components:

$$Y_i | \mathbf{x}_i \sim \text{EDM}(\mu_i, \phi)$$

$$g(\mu_i) = \mathbf{x}_i^T \boldsymbol{\beta}$$

What we know so far

GLMs:

- ▶ Estimating β (and ϕ if needed)
- ▶ Key properties of EDMs

Linear and logistic regression:

- ▶ Model assessment
- ▶ Predictions and prediction error
- ▶ Inference
- ▶ Model selection

Today: Model assumptions and diagnostics for *Poisson* models

Poisson regression

$$Y_i | \mathbf{x}_i \sim \text{Poisson}(\mu_i)$$
$$\log(\mu_i) = \mathbf{x}_i^T \boldsymbol{\beta}$$

Question: What assumptions is this model making? How would we assess those assumptions?


Activity

Use the provided packet to answer the following questions:

- ▶ Does the Poisson regression model appear to be a good fit?
- ▶ If not, which piece of the model do you think is mis-specified?

Rootograms: an overall assessment of model fit

```
topmodels::rootogram(m_art)
```




Rootograms: an overall assessment of model fit

- ▶ horizontal axis: possible values for the response variable
- ▶ vertical axis (square root scale): predicted count for the model (points) and *observed* counts from the data (bars)
- ▶ Assessing whether predicted counts match observed counts
- ▶ Lack of fit could be due to:
 - ▶ Wrong choice of distribution
 - ▶ Additional explanatory variables need to be added to model
 - ▶ Issues with shape assumption

Question: We didn't make a rootogram for binary logistic regression. Why not?

(Randomized) quantile residual plots



What I am looking for in a quantile residual plot?


Examining influential points

Cook's distance:

```
articles[which(cooks.distance(m_art) > 0.1),  
          c("art", "kid5", "phd", "ment")]
```


```
##      art kid5  phd ment  
## 328   1     1 1.78   77  
## 914  16     0 1.74   21  
## 915  19     0 1.86   42
```

(Randomized) quantile residual plots




How could I decide whether deviations in the quantile residual plot are just due to random chance?

Plot lineup




Quantile residual plot

Plot without the two rows with the most number of mentor publications:




QQ plot

If the model is a good fit to the data, randomized quantile residuals should be approximately $N(0, 1)$.



QQ plot

If the model is a good fit to the data, randomized quantile residuals should be approximately $N(0, 1)$.



Summary so far

For the Poisson regression on the articles data, here is my perspective:

- ▶ There are several potential outliers
- ▶ The rootogram shows our predicted counts are somewhat low in the tails of the distribution, compared to the observed counts
- ▶ Quantile residual plots suggest the shape assumption is overall pretty reasonable
- ▶ The QQ plot suggests that the quantile residuals are approximately normal, but variability of the residuals is higher than expected

Question: If this particular model is mis-specified, what do you think is the cause? How could we modify the model?